AMYPO - Policy & Placement Copilot

Manual Test Guide

How to verify the product by hand: what to open, what to type, and what you should see. Written for a reviewer with no prior knowledge of the code.

Applies to: version 0.1.0, build 2026-10-20

Where: http://localhost:3000 after docker compose up

Accounts: admin / staff.cse / s102, password amypo-demo-2026

Time: about 40 minutes for all checks; 10 for the smoke subset

Contents

Section	What it covers
1. Before you start	Stack, accounts, resetting state, how to record a result
2. Coverage map	Which check covers which feature, and where it lives
3. Access and sessions	Six checks: login, roles, sign-out, wrong password
4. Asking questions	Nine checks: citations, trace, cache, abstention, scope
5. Placement	Six checks: eligibility, skill gap, what-if, PDF, staff matching
6. Dashboard and degradation	Seven checks: tiles, Ollama states, scenarios, retrain
7. Knowledge and ingestion	Three checks: document list, re-scan, force re-index
8. Evaluation	Three checks: quick run, full run, role gate
9. Admin and audit	Five checks: users, password reset, audit trail, access
10. API and command line	Five checks: health, smoke suite, model pull, guides
11. Known gaps	What will fail if you test it, and why it is expected
12. Sign-off	One table to record the run

1. Before you start

Bring the stack up from the project root, then wait until the header pill reads **local model** (not **offline mode**) and the Ollama tile says loaded. The first start downloads phi3:mini into the Ollama container, so give it a few minutes on a fresh machine.

```
docker compose up -d --build
```

- Open `http://localhost:3000`. Use a private window so no earlier session survives.
- Use the **Demo accounts** panel on the sign-in page for the three roles; the password is `amypo-demo-2026` for all of them.
- Reset between runs with `docker compose restart backend`; it does not touch the database, the corpus or the model.
- Record PASS or FAIL and the observed text for every id. A check with no observation is not a pass.
- Numbers differ by machine (latency, RAM). Judge the shape of the behaviour, not an exact millisecond.

If the Ollama tile says **model not pulled**, the model is missing from the container's own store: run `docker compose run --rm ollama-pull` (or `make models`). A pull in the host's Ollama app does not count. The card recovers within 30 seconds.

2. Coverage map

Every user-visible capability maps to at least one check. Use this table to decide what a shortened run may skip without leaving a feature unexamined.

Area	Screen / entry point	Checks
Sign-in, sessions, role landing pages	/login, header Sign out	S1-S6
Grounded answers, citations, trace	/ (Chat)	C1-C3
Feedback that feeds retraining	Chat answer footer	C4, D4
Semantic cache	Chat, ask the same question twice	C5
Abstention and closest records	Chat, unanswerable question	C6
Role scoping and staff-only documents	Chat as s102 and staff.cse	C7-C8
Language detection (experimental)	Chat, non-Latin input	C9
History and new chat	Chat toolbar and sidebar	C10
Eligibility, skill gap, what-if	/placement as s102	P1-P3
Eligibility PDF	/placement, Download PDF	P4
Staff candidate matching	/placement as staff.cse	P5
Live resources and model state	/dashboard tiles	D1-D2
Failure-mode scenarios	/dashboard Scenarios	D3
Router retraining	/dashboard Retrain	D4
Cost, latency and quality charts	/dashboard	D5
Decision log	/dashboard, decision table	D6
Corpus and ingestion	/knowledge	K1-K3
Benchmark from the UI	/evaluation	E1-E3
Accounts and password reset	/admin (Users)	A1
Audit trail and denied filter	/audit	A2
Access enforcement	direct URLs and API	A3-A4
Health and readiness	GET /api/v1/health	A5
End-to-end automation	scripts/e2e_smoke.py	X1
Model bootstrap and documents	make models, make guides	X2-X3

3. Access and sessions

ID	What to do	What you should see
S1	Open /dashboard in a private window.	Redirected to /login. After signing in you land back on /dashboard, not the home page.
S2	Demo accounts -> Admin. Sign in.	Lands on Dashboard. Header shows admin / Admin, the local-model pill, and the Users nav item.
S3	Sign out, then sign in as staff.cse.	Lands on Chat. Nav has no Users item. Audit trail is visible.
S4	Sign out, then sign in as s102.	Lands on Chat. No Users item. Placement opens on your own record.

ID	What to do	What you should see
S5	Sign in as admin with password wrong-password.	Inline message: Incorrect username or password. No redirect, no session, no stack trace.
S6	Click Sign out, then press Back.	Login page. The previous page does not render from cache; a reload keeps you signed out.

4. Asking questions

ID	What to do	What you should see
C1	As s102, ask: What is the minimum attendance required for exams?	A direct answer with citation chips. The footer shows a route, latency, confidence, a tier label, a cost of zero and grounded n/n claims.
C2	Click a citation chip such as [S1].	The Sources and trace panel opens on that record. The trace tab lists the real stages: classify, retrieve, gate, generate, verify.
C3	Click Copy under the answer.	Button changes to Copied for about a second; the clipboard holds the answer text.
C4	Click the thumbs-up button in the answer footer.	The icon turns teal and stays pressed. On Dashboard, Retrain now reports confirmed feedback instead of no thumbs feedback yet.
C5	Ask the same question again in the same session.	The route badge shows the cache path and the latency is far below the first answer. The text and citations match.
C6	Ask: What is the Wi-Fi password for the hostel?	An abstention: the answer says it could not find support and asserts nothing. Closest records are listed, and a clarifying question may be offered.
C7	As s102, ask: What is the attendance of S101?	Refused or abstained without any personal row for S101. Sign in as staff.cse and ask again: the cohort answer is allowed.
C8	As staff.cse, ask: What are the internal notes for the placement officer? Then repeat as s102.	Staff get a cited answer from the staff-only document. The student gets no staff-only source and no paraphrase of it.
C9	Ask a question in Devanagari or Tamil script.	A muted line says the script was detected and that the answer is in English. Retrieval still runs; expect a weaker answer.
C10	Click New chat, then reopen the history sidebar.	Empty state returns with starter questions. Previous questions are listed in the sidebar and reopen with their sources.

5. Placement

ID	What to do	What you should see
P1	As s102, open Placement and select a company role.	An Eligibility table with each criterion (CGPA, backlogs, attendance, skills), the observed value and a met/missed verdict.
P2	Scroll to the skill-gap panel.	Missing skills with the level needed, plus a short learning path. No invented certifications.
P3	Raise Improve all skills (levels) by one and click Recompute eligibility.	Opens up / Lost lines and a line like 2 roles eligible now -> 4 after the change, Nothing is saved.
P4	Click Download PDF next to the student id.	A PDF named S102-eligibility.pdf arrives. It contains the criteria table and the learning path, not a blank page.
P5	Sign in as staff.cse, open Placement, select a role.	Eligible for <company> <role> ranks candidates with reasons, and Closest misses lists students one criterion short.
P6	As s102, change the student id field and press Enter.	Either the field is read-only for students or the result still shows only your own record; the server, not the field, decides.

6. Dashboard and degradation

ID	What to do	What you should see
D1	Sign in as admin and open Dashboard.	Five tiles (RAM, CPU, bandwidth, quota, Ollama) refresh every five seconds. Ollama says loaded with a last-reply time.
D2	Stop Ollama, restart the backend, reload the page.	Ollama says unavailable, with a hint naming the start command. Ask a question: it still answers from quoted evidence, labelled extractive. Start Ollama again, restart the backend: loaded returns.
D3	Click each scenario button in turn, ending with reset to stable.	A summary line describes the simulated run. Tiles change (for example bandwidth scenario, quota falling) and the decision log gains rows. Nothing crashes.
D4	Click Retrain.	If C4 was done: a previous -> new agreement line and a confirmation count. If not: no thumbs feedback yet - rate an answer in chat. No error banner.
D5	Compare the three charts against recent activity.	Requests by route, p50/p95 per route and cumulative cost all match the answers you just asked; the cost bar for AMYPO stays far below always-cloud.
D6	Read the decision table under Scenarios.	Each row has a query, the route, complexity, cost and a human-readable reason. Rows accumulate as you ask questions.
D7	Hover the latency bars.	A tooltip gives p50 and p95 in milliseconds for that route. Values are numbers, not placeholders.

7. Knowledge and ingestion

ID	What to do	What you should see
K1	Open Knowledge.	The document count, then a table: title and doc_id, type, access scope, chunk count and indexed time.
K2	Click Re-scan and wait.	A result line reports documents and chunks seen, skipped and re-embedded. The document count does not double.
K3	Click Force re-index, wait, then ask a policy question.	Button reads Ingesting while busy, then the same corpus is rebuilt and the answer is still cited. No duplicate sources appear.

8. Evaluation

ID	What to do	What you should see
E1	As admin, open Evaluation and click Quick run (40).	A completed run with accuracy, groundedness, abstention, hallucination, retrieval and routing figures. The page states it is a live run, not a stored file.
E2	Click Full benchmark, then wait.	The same panel with the full benchmark count. The button returns to Full benchmark; no timeout error appears.
E3	As staff.cse, open /evaluation.	The page either hides the run controls or the run is refused with a clear message; it must never return numbers it did not compute.

9. Admin and audit

ID	What to do	What you should see
A1	As admin, open Users. Filter, disable and re-enable one non-admin account, then reset its password.	Filter narrows the table. Disable flips the button to Enable. Reset needs 8+ characters, then reports success and ends that account's sessions.
A2	Open Audit trail as admin. Tick Denied only.	Allowed, denied and total counts at the top; rows carry time, actor, action, target and result. The filter leaves only denied actions.
A3	As s102, open /admin and /audit directly by URL.	No admin data is shown; the API refuses with 403. The nav never offered those pages.

ID	What to do	What you should see
A4	As staff.cse, open /audit.	The staff view shows cohort queries only - never student record rows, and never the denied-action log.
A5	Open /api/v1/health in a tab.	JSON with status ok and per-component detail for db, index, llm and router; no secret values.

10. API and command line

Run these from the project root while the stack is up.

ID	Command	What you should see
X1	python scripts/e2e_smoke.py	46/46 checks passed, exit code 0. It drives the real HTTP API with the three roles.
X2	docker compose run --rm ollama-pull	model ready: phi3:mini and exit 0; a second run is a no-op, not a re-download.
X3	python scripts/make_guides.py	Three PDFs written under docs/: product guide, user manual and this manual test guide.
X4	curl -s localhost:8000/api/v1/resource-status	JSON with RAM, CPU, bandwidth, quota, ollama_status ok and ollama_detail naming the model.
X5	cd backend && python -m pytest tests -q	The suite passes. Expect roughly ten minutes: the ingest tests rebuild the 8,784-row corpus.

Sample access check without a browser: request `/api/v1/admin/users` with no cookie (expect 401), then with a student session (expect 403). The smoke script performs this for every endpoint and role, so X1 is the cheaper path.

11. Known gaps

These are real limits, not defects to report. If a check contradicts one of them, that is a finding.

- **The cloud leg is simulated.** Quota, price, latency and failure rate are constants; no request leaves the machine. Cost comparisons are therefore model outputs, not bills.
- **Local answers are slow on CPU.** A generated answer takes roughly 10-20 seconds here; cached, tool-backed and extractive answers are far faster. The latency tiles are expected to differ by an order of magnitude.
- **Dense retrieval is effectively off unless sentence-transformers is installed**, so hybrid numbers equal BM25 on a default install. The report says so rather than claiming a gain.
- **Staff have no PDF control in the UI.** The API supports staff exports, but the button is shown only for a student's own report.
- **Vague questions abstain.** 'What is the attendance policy?' is below the relevance floor; naming the requirement gets an answer. This is the gate working, not a retrieval crash.
- **The suite is slow by design.** Ingest tests rebuild the corpus, so the full run is minutes, not seconds. Run it once, not per check.

12. Sign-off

Record the run against the build date on the cover. A shortened run is acceptable if the skipped ids are named; an unrecorded observation is not.

Section	Checks	Passed	Notes
3. Access and sessions	S1-S6	___ / 6	
4. Asking questions	C1-C10	___ / 10	
5. Placement	P1-P6	___ / 6	

Section	Checks	Passed	Notes
6. Dashboard	D1-D7	___ / 7	
7. Knowledge	K1-K3	___ / 3	
8. Evaluation	E1-E3	___ / 3	
9. Admin and audit	A1-A5	___ / 5	
10. Command line	X1-X5	___ / 5	
Total	45	___ / 45	

Tester: _____ Date: _____ Build: 2026-10-20

Signature: _____ Environment: laptop / demo machine / other